

Deep Learning Course – L 7 - Detailed Notes

Gradient Descent and Learning Rate in Neural Networks

Gradient Descent is one of the most essential optimization algorithms used in machine learning and deep learning. In this lecture, we'll explore **how gradient descent works** and the **critical role of the learning rate** in optimizing a model's parameters.

What Is Gradient Descent?

Forward Propagation and Backward Propagation in Neural Networks

1. Recap of Previous Concepts

Before diving into forward and backward propagation, let's briefly revisit what we've covered:

Logistic Regression

- Logistic regression is used for binary classification.
- It forms the foundation of a simple neural network (a single-layer perceptron).
- The output is passed through a **sigmoid activation function** to predict probabilities.

Weights and Bias

- **Weights (W):** Control the strength of the connection between neurons.
- **Bias (b):** Allows the model to shift the activation function to fit the data better.
- Both are learnable parameters.

Loss Function

- Measures the difference between actual output (label y) and predicted output (\hat{y}).
- Common loss for classification: **Binary Cross-Entropy Loss**.

- Helps assess how well the model is performing.

✓ Gradient Descent

- Optimization algorithm used to minimize the loss.
- Updates weights and biases in the opposite direction of the gradient (steepest descent).
- Formula:

$$W = W - \alpha \cdot \frac{\partial \text{Loss}}{\partial W}$$

where α is the learning rate.

🔄 2. Forward and Backward Propagation

These are the **core components of training** in neural networks.

◆ Forward Propagation

What is it?

- The process of **passing input data through the network** to get a prediction.

Steps:

1. **Input layer:** Takes raw input x .
2. **Hidden layers:** Compute weighted sum \rightarrow apply activation functions.
3. **Output layer:** Produces the prediction \hat{y} .
4. **Loss is calculated** by comparing \hat{y} with actual output y .

Example:

For one hidden layer:

$$\begin{aligned} z[1] &= W[1]x + b[1] \\ a[1] &= \text{ReLU}(z[1]) \\ z[2] &= W[2]a[1] + b[2] \\ \hat{y} &= \sigma(z[2]) \end{aligned}$$

Output:

- We get the predicted value \hat{y} .

- Next, we compute the loss.
-

◆ Backward Propagation

What is it?

- The process of **updating the weights and biases** to minimize the loss using derivatives.

Goal:

- Understand how changing each parameter (weight/bias) affects the loss.

Uses:

- **Chain Rule** to compute gradients from output layer back to input.
-

🧩 3. Toy Example: Step-by-Step

Let's understand forward and backward propagation using a very basic mathematical function:

Function:

$$j = 3 \times ((a + b \times c))$$

Step 1: Define Operations

- Let $u = b \times c$
 - Let $v = a + u$
 - Then $j = 3 \times v$
-

🌀 Forward Propagation

Let's assign values:

- $a = 1, b = 2, c = 3$

Now compute:

1. $u=2 \times 3=6$ $u = 2 \times 3 = 6$ $u=2 \times 3=6$
2. $v=1+6=7$ $v = 1 + 6 = 7$ $v=1+6=7$
3. $j=3 \times 7=21$ $j = 3 \times 7 = 21$ $j=3 \times 7=21$

So, the output is **21**.

Backward Propagation

Now we want to compute **how sensitive the output j is to each variable (a, b, c)**.

We use **derivatives**:

Step 1: Derivative of j w.r.t v :

$$\frac{\partial j}{\partial v} = 3 \frac{\partial j}{\partial v} = 3 \frac{\partial j}{\partial v} = 3$$

Step 2: Derivative of v w.r.t a :

$$\frac{\partial v}{\partial a} = 1 \Rightarrow \frac{\partial j}{\partial a} = \frac{\partial j}{\partial v} \cdot \frac{\partial v}{\partial a} = 3 \cdot 1 = 3 \quad \frac{\partial j}{\partial a} = \frac{\partial j}{\partial v} \cdot \frac{\partial v}{\partial a} = 3 \cdot 1 = 3 \quad \frac{\partial j}{\partial a} = \frac{\partial j}{\partial v} \cdot \frac{\partial v}{\partial a} = 3 \cdot 1 = 3$$

Step 3: Derivative of v w.r.t u :

$$\frac{\partial v}{\partial u} = 1 \quad \frac{\partial v}{\partial u} = 1 \quad \frac{\partial v}{\partial u} = 1$$

Step 4: Derivative of u w.r.t b and c :

$$\frac{\partial u}{\partial b} = c = 3, \frac{\partial u}{\partial c} = b = 2 \quad \frac{\partial u}{\partial b} = c = 3, \quad \frac{\partial u}{\partial c} = b = 2$$

Now use chain rule:

$$\frac{\partial j}{\partial b} = \frac{\partial j}{\partial v} \cdot \frac{\partial v}{\partial u} \cdot \frac{\partial u}{\partial b} = 3 \cdot 1 \cdot 3 = 9 \quad \frac{\partial j}{\partial b} = \frac{\partial j}{\partial v} \cdot \frac{\partial v}{\partial u} \cdot \frac{\partial u}{\partial b} = 3 \cdot 1 \cdot 3 = 9$$

$$\frac{\partial j}{\partial c} = \frac{\partial j}{\partial v} \cdot \frac{\partial v}{\partial u} \cdot \frac{\partial u}{\partial c} = 3 \cdot 1 \cdot 2 = 6 \quad \frac{\partial j}{\partial c} = \frac{\partial j}{\partial v} \cdot \frac{\partial v}{\partial u} \cdot \frac{\partial u}{\partial c} = 3 \cdot 1 \cdot 2 = 6$$

4. Chain Rule in Backpropagation

The **chain rule** is the backbone of backward propagation. It allows us to calculate:

$$\frac{\partial j}{\partial x} = \frac{\partial j}{\partial y} \cdot \frac{\partial y}{\partial x} \quad \frac{\partial j}{\partial x} = \frac{\partial j}{\partial y} \cdot \frac{\partial y}{\partial x}$$

In complex neural networks, we stack multiple functions together, so we apply the chain rule repeatedly through layers to calculate gradients.

5. Why Is Backpropagation Important?

- It gives us the **direction** to move weights to reduce the loss.
 - Used along with gradient descent to **train** the neural network.
 - Efficient for **deep networks** with many layers.
-

6. Summary of Key Points

| Concept | Description |
|----------------------|--|
| Forward Propagation | Pass input through network to get output |
| Loss Function | Measures error between prediction and actual output |
| Backward Propagation | Calculates gradients of weights & biases |
| Chain Rule | Helps in computing gradients across multiple layers |
| Optimization | Gradients used to update parameters and reduce error |

What's Next?

In the next session, we'll study **activation functions** (like Sigmoid, ReLU, Tanh), why they are essential, and how they add non-linearity to neural networks.

.